

Brian D'Alelio DSE5002 Project 1

The CEO wants to understand how much to pay a data scientist, and potentially a team of data scientists in the future, both in the United States and Internationally. She understands that salaries vary widely across the world which creates uncertainty around this decision. The job landscape is seeing salaries increasing due to the great recession making the market highly competitive. My task is to analyze data science salaries and recommend a salary range. I'll present my recommendations visually in a Powerpoint presentation.

What should the company offer as a competitive salary for a data scientist? How do data science salaries differ around the world, and what pay range would attract strong talent? Given global salary trends and rising competition, what is an appropriate compensation package for a full-time data scientist? Should the analysis account for experience level?

```
In [11]: import pandas as pd
df = pd.read_csv("r project data-1-1.csv")
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 607 entries, 0 to 606
Data columns (total 12 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Unnamed: 0        607 non-null    int64  
 1   work_year         607 non-null    int64  
 2   experience_level 607 non-null    object  
 3   employment_type   607 non-null    object  
 4   job_title          607 non-null    object  
 5   salary             607 non-null    int64  
 6   salary_currency   607 non-null    object  
 7   salary_in_usd     607 non-null    int64  
 8   employee_residence 607 non-null    object  
 9   remote_ratio       607 non-null    int64  
 10  company_location   607 non-null    object  
 11  company_size       607 non-null    object  
dtypes: int64(5), object(7)
memory usage: 57.0+ KB
```

```
In [12]: df['experience_level']=pd.Categorical(df.experience_level)
df['employment_type']=pd.Categorical(df.employment_type)
df['job_title']=pd.Categorical(df.job_title)
df['salary_currency']=pd.Categorical(df.salary_currency)
df['employee_residence']=pd.Categorical(df.employee_residence)
df['company_location']=pd.Categorical(df.company_location)
df['company_size']=pd.Categorical(df.company_size)
df['remote_ratio']=pd.Categorical(df.remote_ratio)

df.dtypes
```

```
Out[12]: Unnamed: 0          int64
work_year           int64
experience_level   category
employment_type    category
job_title          category
salary              int64
salary_currency    category
salary_in_usd      int64
employee_residence category
remote_ratio       category
company_location   category
company_size       category
dtype: object
```

```
In [13]: df.describe(include='all')
```

	Unnamed: 0	work_year	experience_level	employment_type	job_title	salar
count	607.000000	607.000000		607	607	607 6.070000e+C
unique	Nan	Nan		4	4	50 N
top	Nan	Nan	SE	FT	Data Scientist	N
freq	Nan	Nan	280	588	143	N
mean	303.000000	2021.405272		Nan	Nan	3.240001e+C
std	175.370085	0.692133		Nan	Nan	1.544357e+C
min	0.000000	2020.000000		Nan	Nan	4.000000e+C
25%	151.500000	2021.000000		Nan	Nan	7.000000e+C
50%	303.000000	2022.000000		Nan	Nan	1.150000e+C
75%	454.500000	2022.000000		Nan	Nan	1.650000e+C
max	606.000000	2022.000000		Nan	Nan	3.040000e+C

```
In [14]: df['experience_level'].value_counts()
```

```
Out[14]: experience_level
SE    280
MI    213
EN    88
EX    26
Name: count, dtype: int64
```

```
In [15]: df['employment_type'].value_counts()
```

```
Out[15]: employment_type  
FT      588  
PT      10  
CT       5  
FL       4  
Name: count, dtype: int64
```

```
In [16]: df['job_title'].value_counts().head(20)
```

```
Out[16]: job_title  
Data Scientist          143  
Data Engineer           132  
Data Analyst            97  
Machine Learning Engineer 41  
Research Scientist      16  
Data Science Manager    12  
Data Architect          11  
Big Data Engineer       8  
Machine Learning Scientist 8  
Data Analytics Manager   7  
Data Science Consultant   7  
Director of Data Science 7  
Principal Data Scientist 7  
AI Scientist             7  
Computer Vision Engineer 6  
BI Data Analyst          6  
ML Engineer              6  
Lead Data Engineer        6  
Business Data Analyst     5  
Applied Data Scientist    5  
Name: count, dtype: int64
```

```
In [17]: df['employee_residence'].value_counts().head(20)
```

```
Out[17]: employee_residence
US      332
GB       44
IN       30
CA       29
DE       25
FR       18
ES       15
GR       13
JP        7
PK        6
BR        6
PT        6
NL        5
IT        4
PL        4
RU        4
TR        3
AU        3
VN        3
AT        3
Name: count, dtype: int64
```

```
In [18]: df['company_location'].value_counts().head(20)
```

```
Out[18]: company_location
US      355
GB       47
CA       30
DE       28
IN       24
FR       15
ES       14
GR       11
JP        6
AT        4
PL        4
NL        4
PT        4
AE        3
BR        3
DK        3
AU        3
LU        3
MX        3
TR        3
Name: count, dtype: int64
```

```
In [19]: df['company_size'].value_counts()
```

```
Out[19]: company_size
M      326
L      198
S       83
Name: count, dtype: int64
```

```
In [20]: df['remote_ratio'].value_counts()
```

```
Out[20]: remote_ratio
100    381
0     127
50     99
Name: count, dtype: int64
```

```
In [21]: df['salary_in_usd'].describe()
```

```
Out[21]: count      607.000000
mean     112297.869852
std      70957.259411
min      2859.000000
25%     62726.000000
50%     101570.000000
75%     150000.000000
max     600000.000000
Name: salary_in_usd, dtype: float64
```

```
In [22]: df['salary'].describe()
```

```
Out[22]: count     6.070000e+02
mean      3.240001e+05
std       1.544357e+06
min      4.000000e+03
25%     7.000000e+04
50%     1.150000e+05
75%     1.650000e+05
max     3.040000e+07
Name: salary, dtype: float64
```

```
In [23]: df['salary_currency'].value_counts()
```

```
Out[23]: salary_currency
USD    398
EUR    95
GBP    44
INR    27
CAD    18
JPY     3
PLN     3
TRY     3
AUD     2
BRL     2
HUF     2
CNY     2
DKK     2
MXN     2
SGD     2
CHF     1
CLP     1
Name: count, dtype: int64
```

```
In [24]: df['work_year'].value_counts()
```

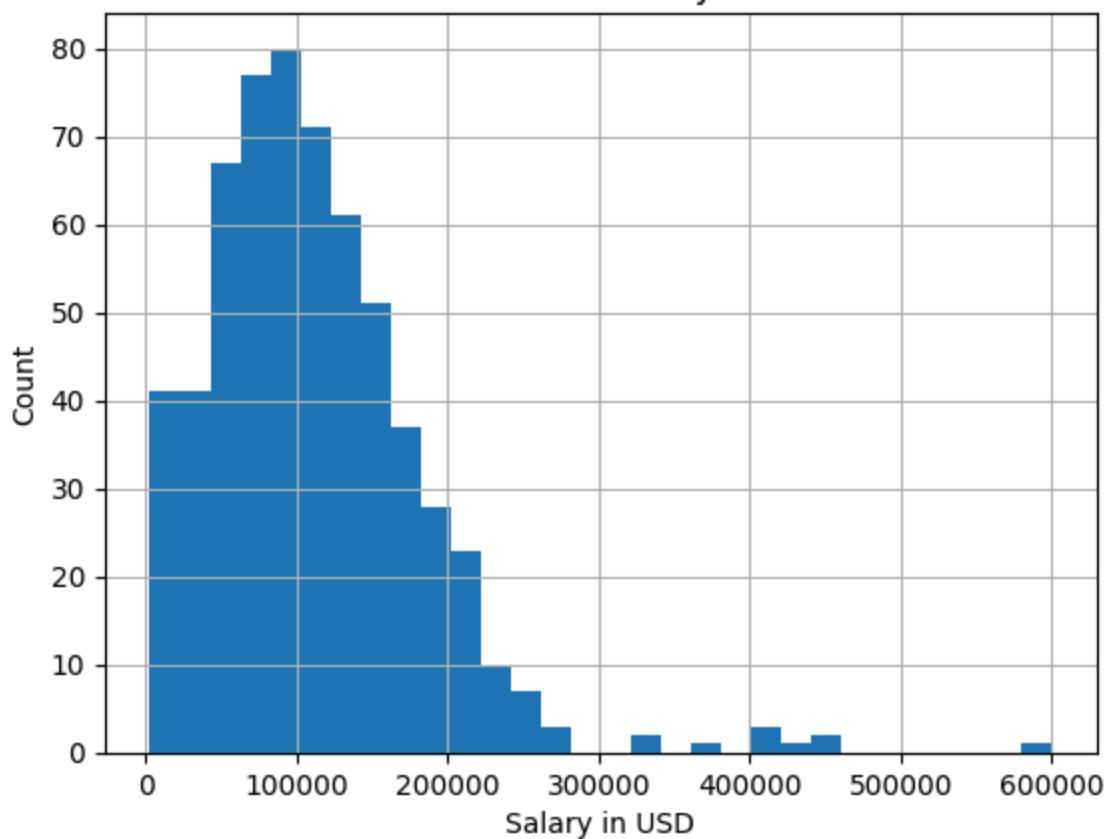
```
Out[24]: work_year
2022    318
2021    217
2020     72
Name: count, dtype: int64
```

```
In [ ]: work_year
        Values range from 2020-2022
        Most records are recent as Median is 2022
experience_level
        4 unique levels - EN, MI, SE, EX
employment_type
        4 categories. Full time (FT) most common by a wide margin
job_title
        50 unique titles. Most common is "Data Scientist".
        Data Engineer, Data Analyst, Machine Learning Engineer also significantly represent
salary
        Salary_in_usd will be cleaner to work with as this is showing salaries in many
salary_currency
        most salaries are denominated in USD. The Euro and Great British Pound also have
salary_in_usd
        Large SD and min max disparity due to differences in geography, experience level
        Median is lower than the mean. This suggests a small number of high earners draw
employee_residence
        There are many countries represented in the data, but the United States is the
remote_ratio
        Fully remote roles are the most common, aligning with a global shift in this dimension
company_location
        company location looks similar to employee residence. This makes sense as employees
        employer of the same country. The differences we do see here indicate cross-border
        at least somewhat common.
company_size
        Three categories - Large, Medium, and Small. Medium and large companies make up
        of the data set, with Medium leading. Small companies appear less often, but still
        meaningful share.
```

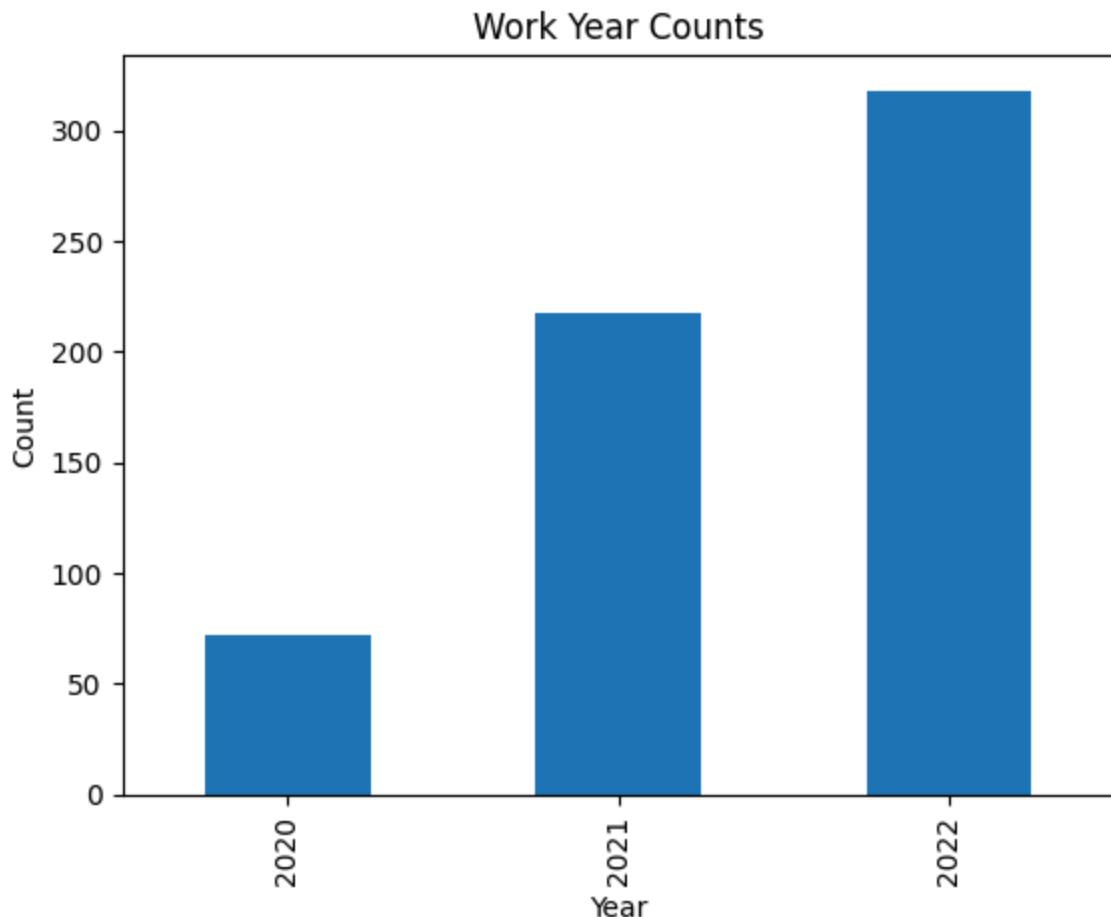
```
In [25]: import matplotlib.pyplot as plt
```

```
In [26]: df['salary_in_usd'].hist(bins=30)
plt.title("Distribution of Salary in USD")
plt.xlabel("Salary in USD")
plt.ylabel("Count")
plt.show()
```

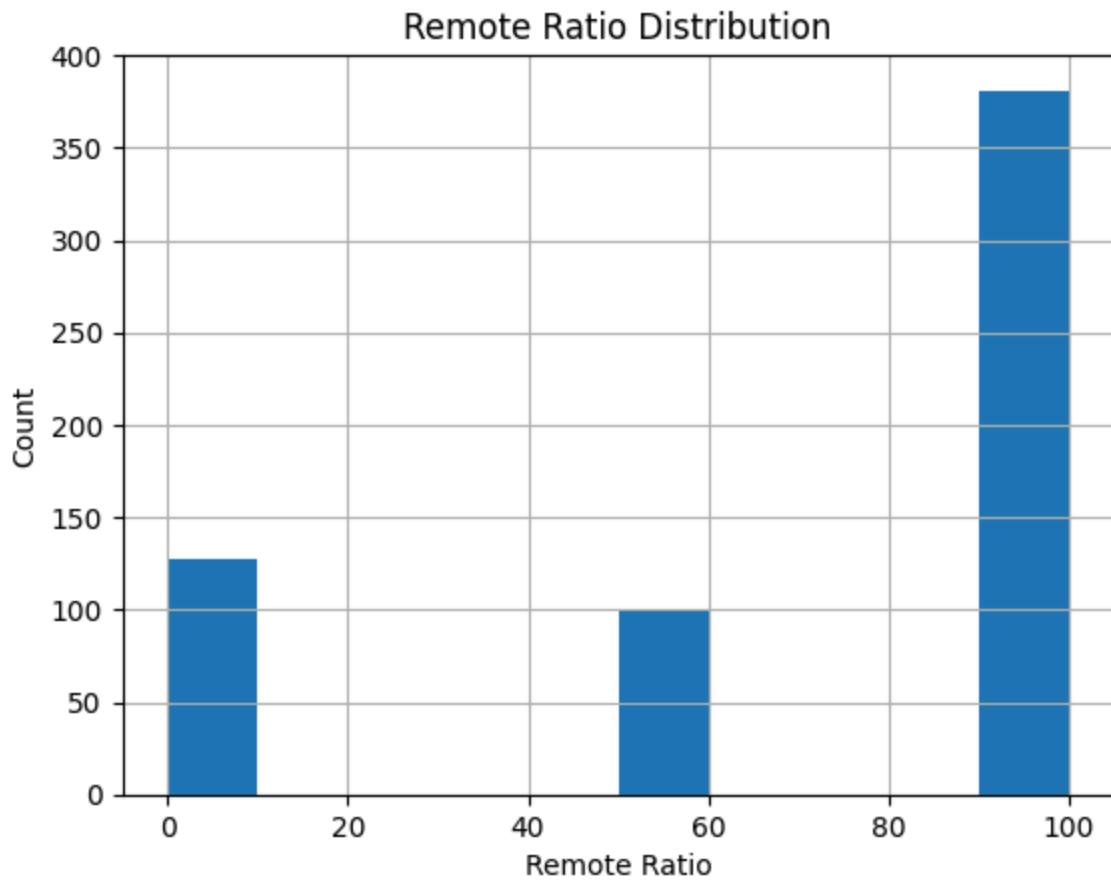
Distribution of Salary in USD



```
In [27]: df['work_year'].value_counts().sort_index().plot(kind='bar')
plt.title("Work Year Counts")
plt.xlabel("Year")
plt.ylabel("Count")
plt.show()
```

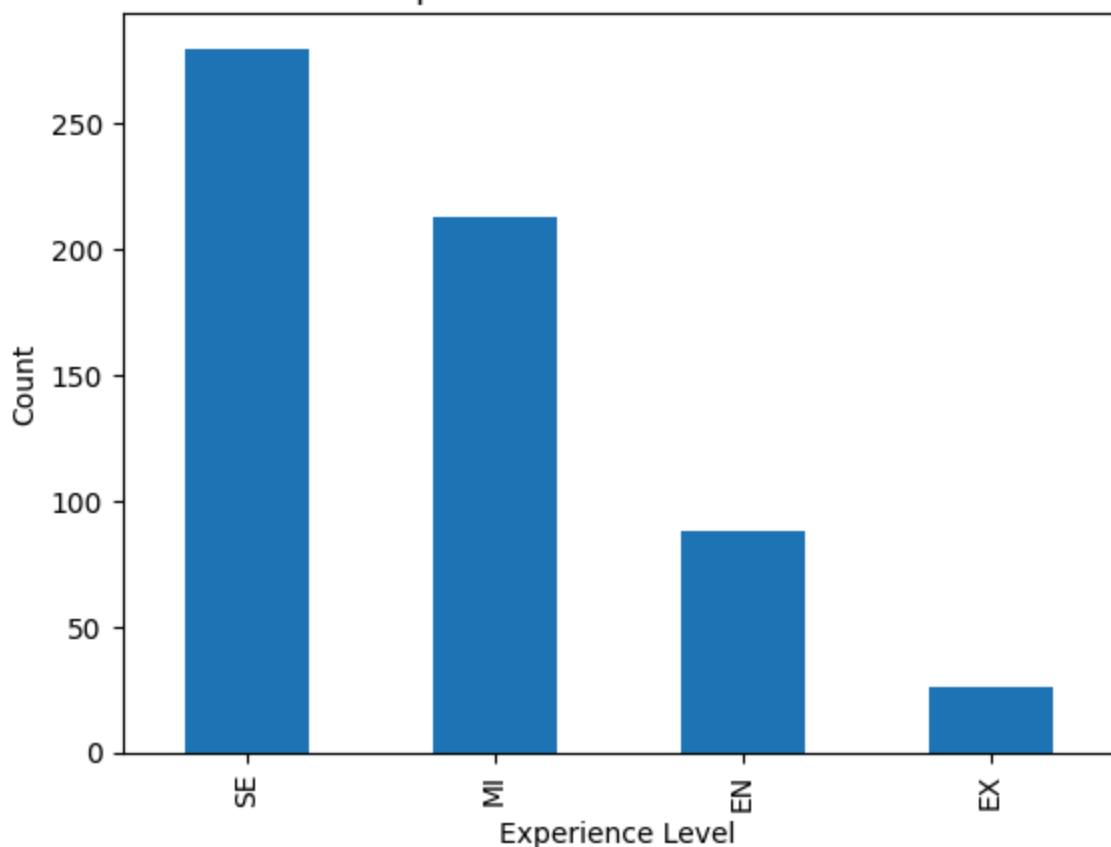


```
In [28]: df['remote_ratio'].hist()  
plt.title("Remote Ratio Distribution")  
plt.xlabel("Remote Ratio")  
plt.ylabel("Count")  
plt.show()
```

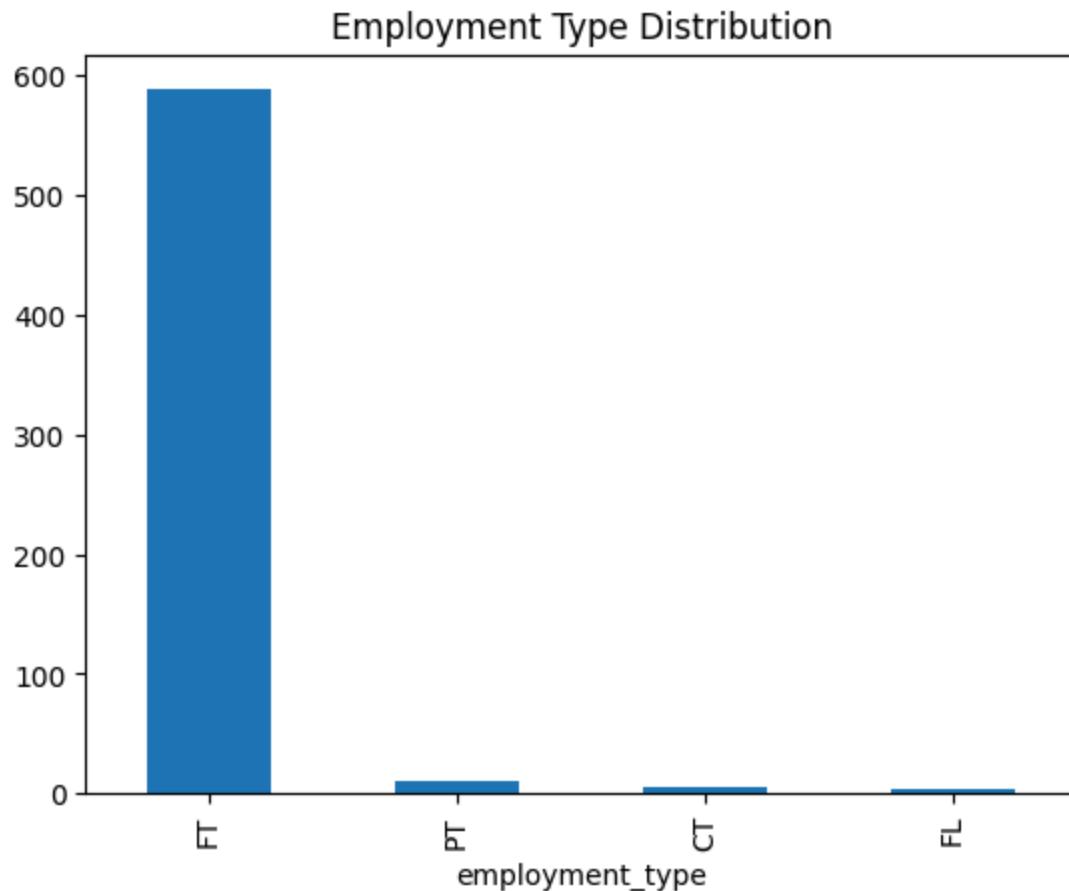


```
In [29]: df['experience_level'].value_counts().plot(kind='bar')
plt.title("Experience Level Distribution")
plt.xlabel("Experience Level")
plt.ylabel("Count")
plt.show()
```

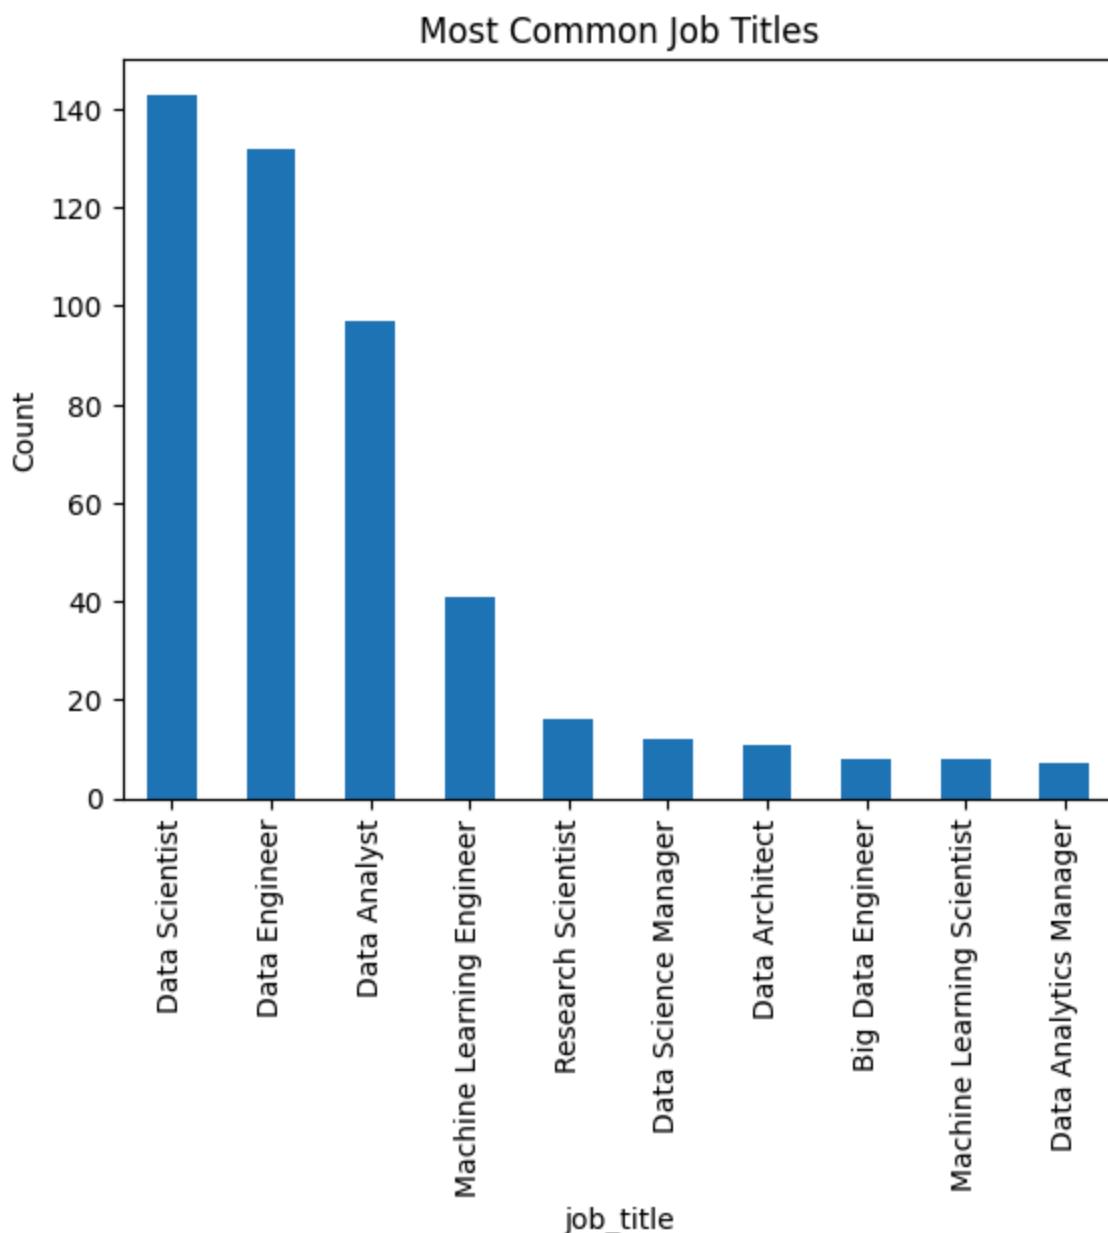
Experience Level Distribution



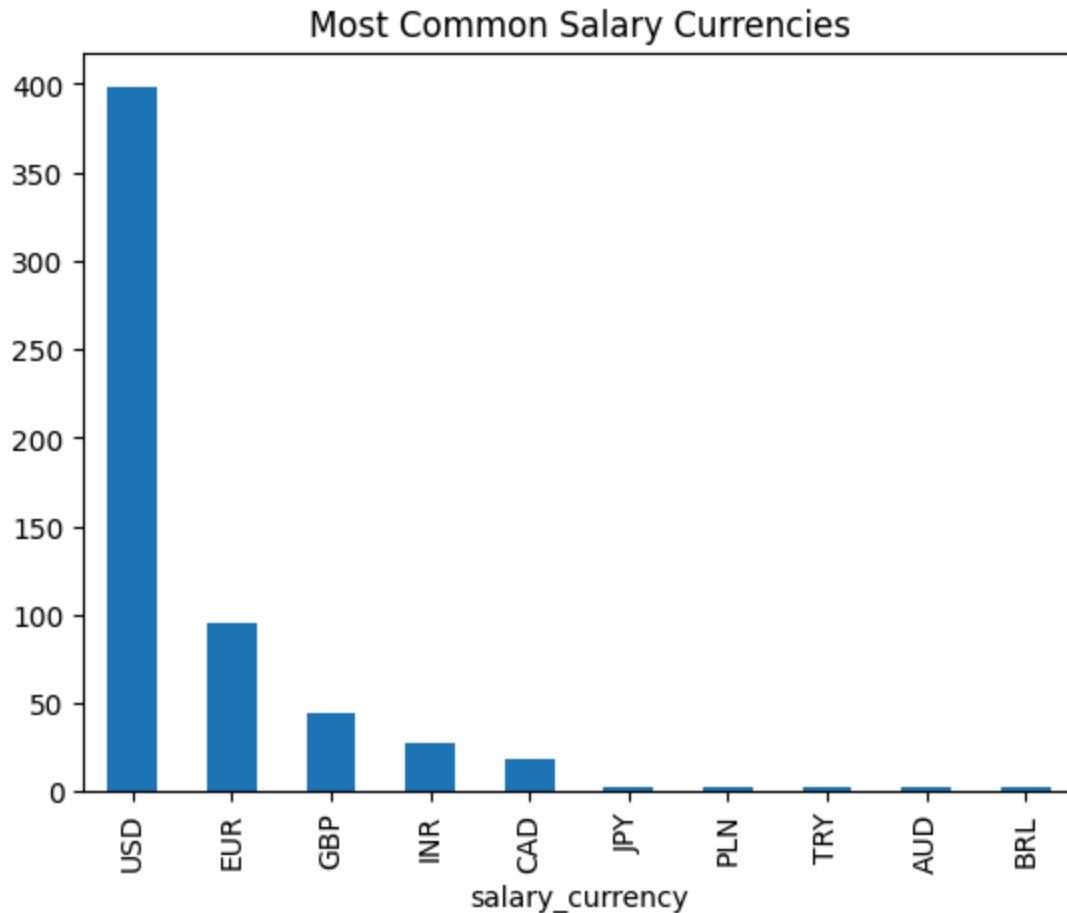
```
In [30]: df['employment_type'].value_counts().plot(kind='bar')
plt.title("Employment Type Distribution")
plt.show()
```



```
In [31]: df['job_title'].value_counts().head(10).plot(kind='bar')
plt.title("Most Common Job Titles")
plt.ylabel("Count")
plt.show()
```

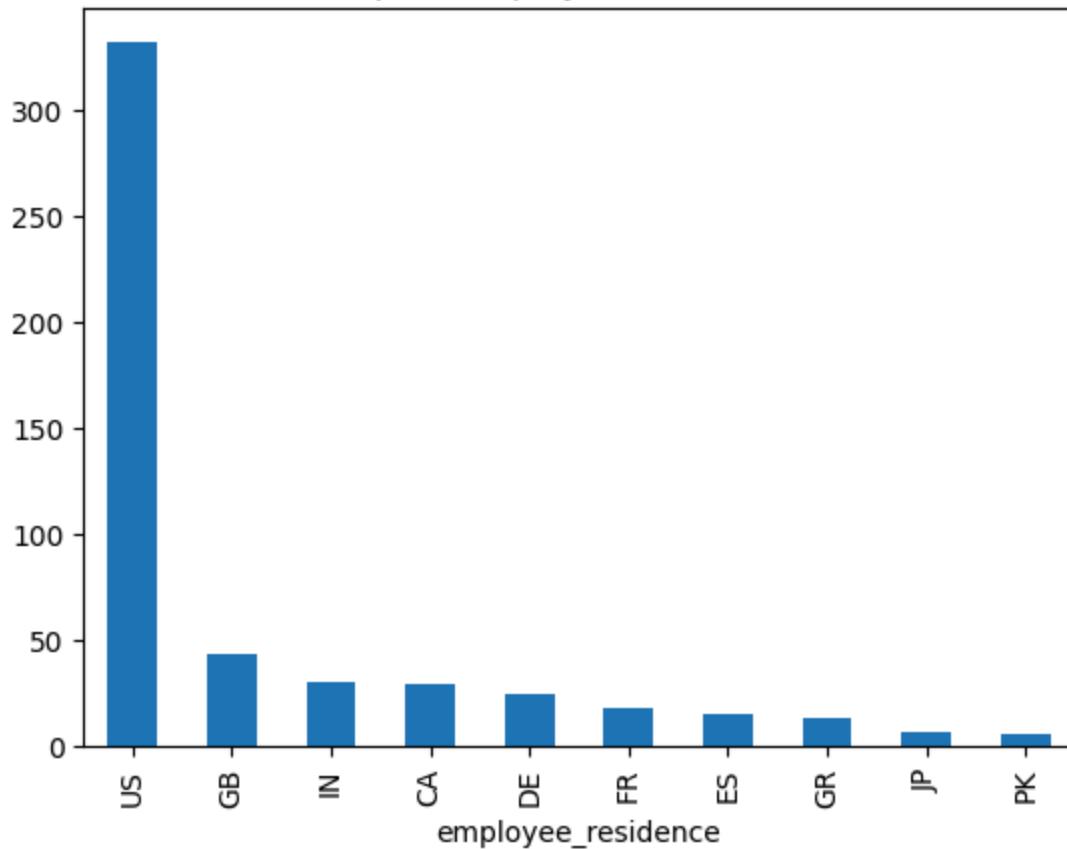


```
In [32]: df['salary_currency'].value_counts().head(10).plot(kind='bar')
plt.title("Most Common Salary Currencies")
plt.show()
```



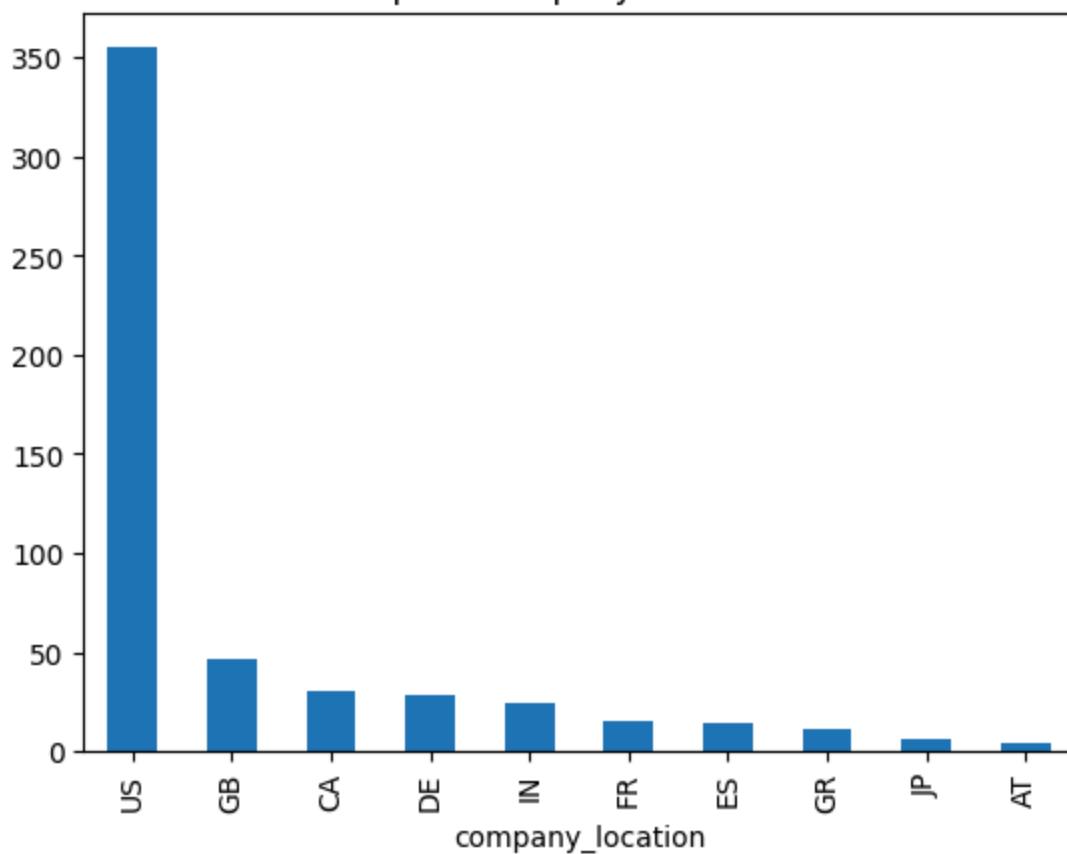
```
In [33]: df['employee_residence'].value_counts().head(10).plot(kind='bar')
plt.title("Top 10 Employee Residences")
plt.show()
```

Top 10 Employee Residences



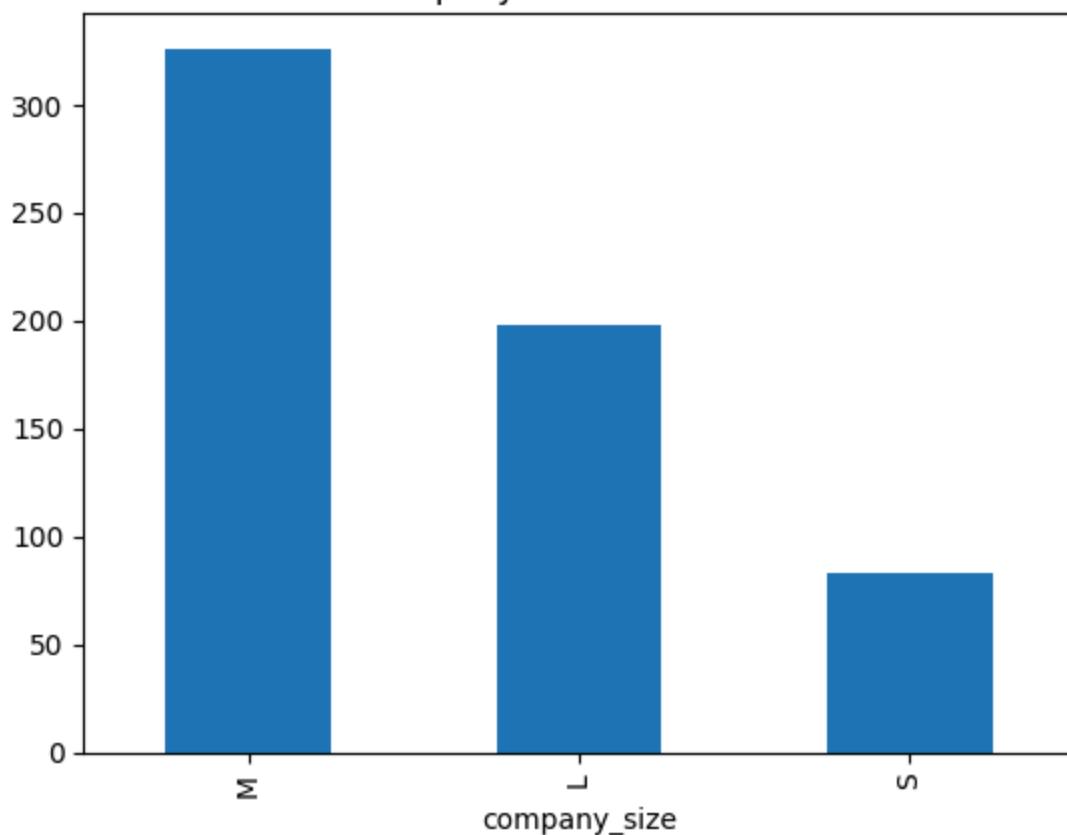
```
In [34]: df['company_location'].value_counts().head(10).plot(kind='bar')
plt.title("Top 10 Company Locations")
plt.show()
```

Top 10 Company Locations



```
In [35]: df['company_size'].value_counts().plot(kind='bar')
plt.title("Company Size Distribution")
plt.show()
```

Company Size Distribution



```
In [36]: exp_map = {
    "EN": "Entry-level",
    "MI": "Mid-level",
    "SE": "Senior-level",
    "EX": "Executive"
}

exp_level = (
    df.groupby("experience_level")["salary_in_usd"]
    .mean()
    .round(0)
    .reset_index()
)

exp_level["experience_description"] = exp_level["experience_level"].map(exp_map)

exp_level = exp_level.sort_values("salary_in_usd", ascending=False)
exp_level
```

C:\Users\Brian\AppData\Local\Temp\ipykernel_22240\891921729.py:9: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.
df.groupby("experience_level")["salary_in_usd"]

	experience_level	salary_in_usd	experience_description
1	EX	199392.0	Executive
3	SE	138617.0	Senior-level
2	MI	87996.0	Mid-level
0	EN	61643.0	Entry-level

```
In [37]: import pycountry

def country_name(code):
    try:
        return pycountry.countries.get(alpha_2=code).name
    except:
        mapping = {"GR": "Greece", "PT": "Portugal"}
        return mapping.get(code, None)

emp_residence = (
    df.groupby("employee_residence")["salary_in_usd"]
        .agg(mean_salary_usd="mean", median_salary_usd="median")
        .round(0)
        .reset_index()
)

emp_residence["country_name"] = emp_residence["employee_residence"].apply(country_n
emp_residence.sort_values("mean_salary_usd", ascending=False, inplace=True)
emp_residence.head(10)
```

C:\Users\Brian\AppData\Local\Temp\ipykernel_22240\626899275.py:11: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.
df.groupby("employee_residence")["salary_in_usd"]

	employee_residence	mean_salary_usd	median_salary_usd	country_name
38	MY	200000.0	200000.0	Malaysia
45	PR	160000.0	160000.0	Puerto Rico
55	US	149194.0	138475.0	United States
41	NZ	125000.0	125000.0	New Zealand
9	CH	122346.0	122346.0	Switzerland
3	AU	108043.0	87425.0	Australia
49	RU	105750.0	72500.0	Russian Federation
50	SG	104176.0	104176.0	Singapore
32	JP	103538.0	74000.0	Japan
31	JE	100000.0	100000.0	Jersey

```
In [38]: job_title = (
    df.groupby("job_title")
    .agg(
        count=("job_title", "size"),
        mean_salary_usd=("salary_in_usd", "mean"),
        median_salary_usd=("salary_in_usd", "median")
    )
    .sort_values("count", ascending=False)
    .head(15)
    .round(0)
    .reset_index()
)

job_title
```

C:\Users\Brian\AppData\Local\Temp\ipykernel_22240\1116427211.py:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
df.groupby("job_title")
```

	job_title	count	mean_salary_usd	median_salary_usd
0	Data Scientist	143	108188.0	103691.0
1	Data Engineer	132	112725.0	105500.0
2	Data Analyst	97	92893.0	90320.0
3	Machine Learning Engineer	41	104880.0	87932.0
4	Research Scientist	16	109020.0	76264.0
5	Data Science Manager	12	158328.0	155750.0
6	Data Architect	11	177874.0	180000.0
7	Big Data Engineer	8	51974.0	41306.0
8	Machine Learning Scientist	8	158412.0	156500.0
9	Data Analytics Manager	7	127134.0	120000.0
10	Data Science Consultant	7	69421.0	76833.0
11	Director of Data Science	7	195074.0	168000.0
12	Principal Data Scientist	7	215242.0	173762.0
13	AI Scientist	7	66136.0	45896.0
14	Computer Vision Engineer	6	44419.0	26304.0

```
In [39]: import seaborn as sns

exp_level = df.groupby("experience_level")["salary_in_usd"].agg(["mean", "median"])

exp_level["experience_level"] = exp_level["experience_level"].map({
```

```

    "EN": "Entry-level",
    "MI": "Mid-level",
    "SE": "Senior-level",
    "EX": "Executive-level"
})

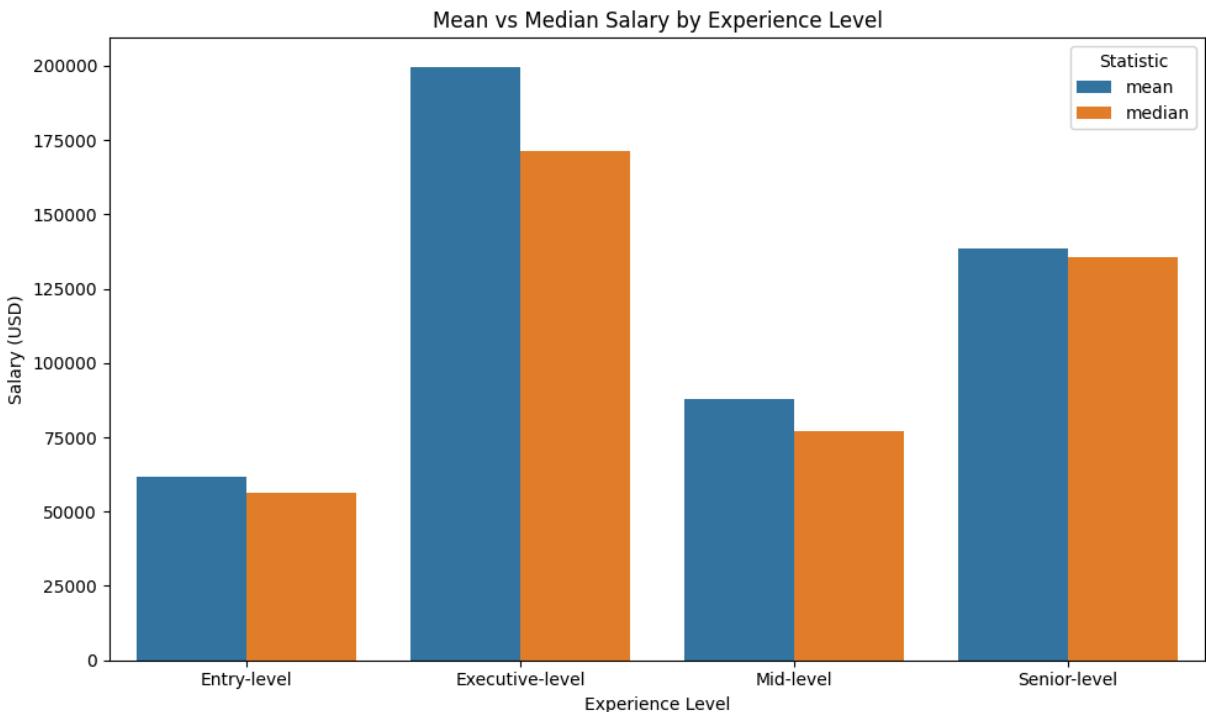
exp_level_melted = exp_level.melt(id_vars="experience_level", value_vars=["mean", "median"], var_name="Statistic", value_name="Salary")

plt.figure(figsize=(10,6))
sns.barplot(data=exp_level_melted, x="experience_level", y="Salary", hue="Statistic")
plt.xlabel("Experience Level")
plt.ylabel("Salary (USD)")
plt.title("Mean vs Median Salary by Experience Level")
plt.tight_layout()
plt.show()

```

C:\Users\Brian\AppData\Local\Temp\ipykernel_22240\2736388065.py:3: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
exp_level = df.groupby("experience_level")["salary_in_usd"].agg(["mean", "median"]).reset_index()
```



```
In [44]: top10 = (
    emp_residence
    .dropna(subset=["country_name"])
    .sort_values("mean_salary_usd", ascending=False)
    .head(10)
)

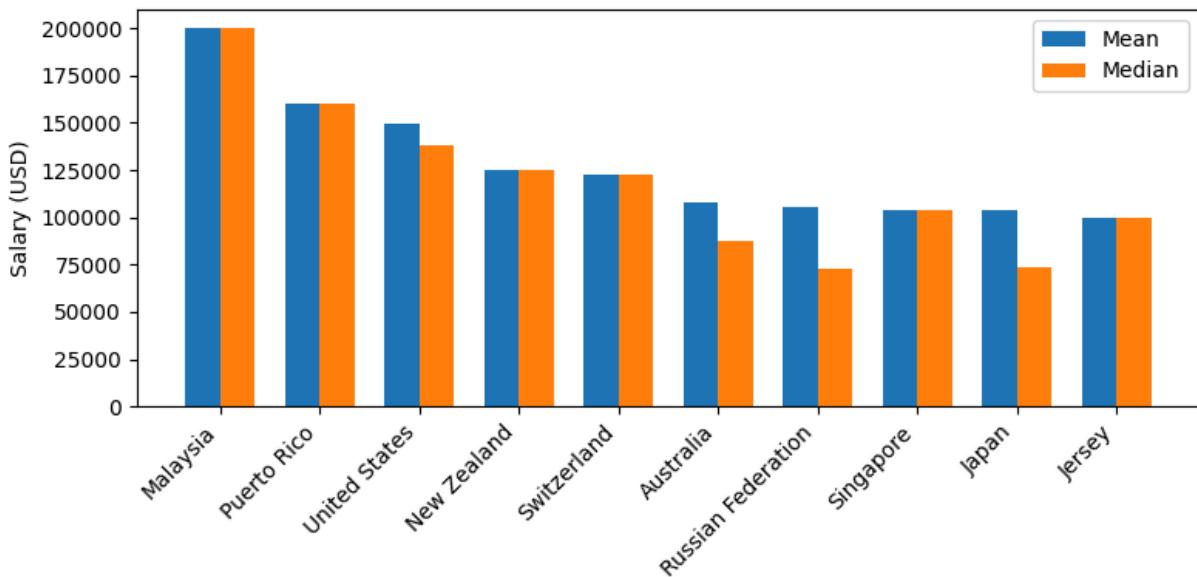
plt.figure(figsize=(8,4))
x = range(len(top10))
width = 0.35
```

```

plt.bar([i - width/2 for i in x], top10["mean_salary_usd"], width, label="Mean")
plt.bar([i + width/2 for i in x], top10["median_salary_usd"], width, label="Median")

plt.xticks(list(x), top10["country_name"], rotation=45, ha="right")
plt.ylabel("Salary (USD)")
plt.legend()
plt.tight_layout()
plt.show()

```



```

In [46]: job_title = (
    df.groupby("job_title")["salary_in_usd"]
        .agg(mean_salary_usd="mean", median_salary_usd="median", count="size")
        .reset_index()
)

top_jobs = job_title.sort_values("count", ascending=False).head(10)

plt.figure(figsize=(8,5))

y = range(len(top_jobs))
width = 0.35

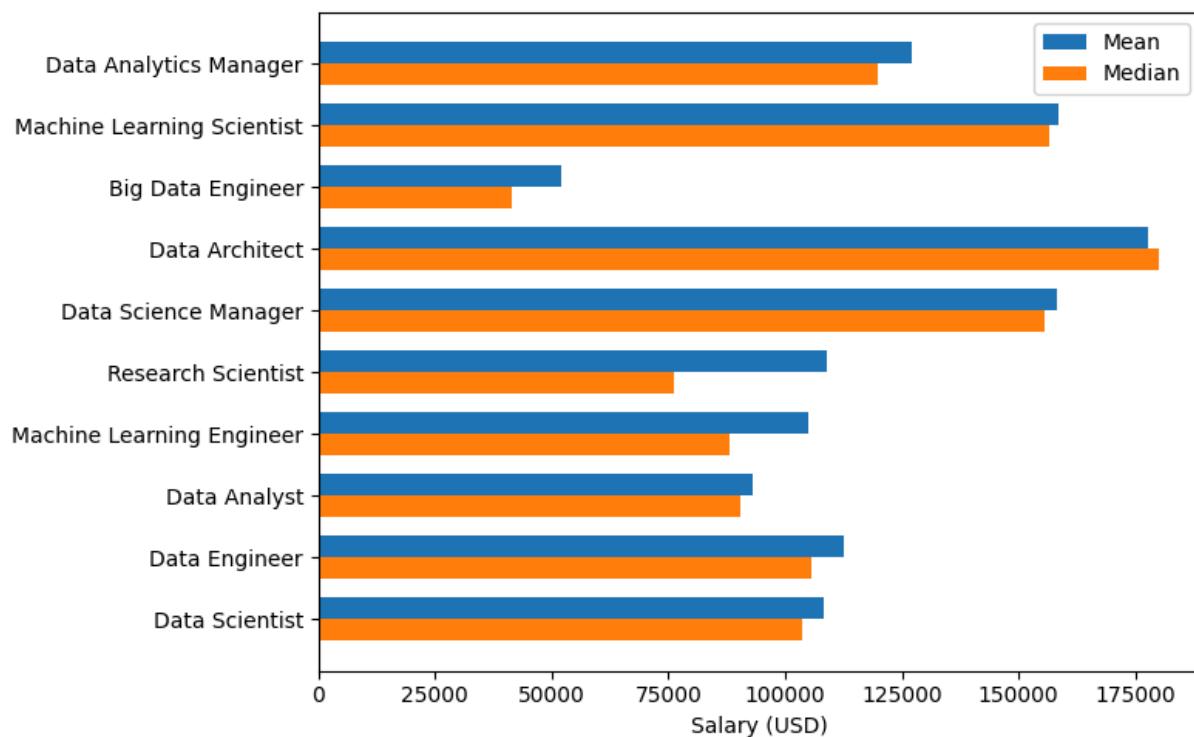
plt.barh([i + width/2 for i in y], top_jobs["mean_salary_usd"], height=width, label="Mean")
plt.barh([i - width/2 for i in y], top_jobs["median_salary_usd"], height=width, label="Median")

plt.yticks(y, top_jobs["job_title"])
plt.xlabel("Salary (USD)")
plt.legend()
plt.tight_layout()
plt.show()

```

C:\Users\Brian\AppData\Local\Temp\ipykernel_22240\410625151.py:2: FutureWarning: The default of observed=False is deprecated and will be changed to True in a future version of pandas. Pass observed=False to retain current behavior or observed=True to adopt the future default and silence this warning.

```
df.groupby("job_title")["salary_in_usd"]
```



Summary:

This analysis explores global data science salaries to identify a competitive pay range for hiring a full-time data scientist. Salaries increase significantly by experience level, vary by employee location (with U.S. workers earning among the highest), and differ across specific job titles. By combining these factors, we can identify salary expectations that will position the company to attract strong candidates in a competitive, rapidly growing market.